


# The editor your agent wishes it had.

A grep-fast, zero-dependency CLI  
*built by an AI agent*

*for the way 1,100+ measured agent sessions actually read & edit code.*



CU · THE CURSOR, ALIVE

tarn is a tiny terminal editor and a structural command-line toolkit for AI coding agents: orient in a repo, jump to a symbol's definition or its uses, read one function, and edit by unit-of-meaning — all deterministic, --json-chainable, with meaningful exit codes. Pure Rust, **no crate dependencies**. And on large-file counting it is at parity with — or faster than — ripgrep, the fastest search tool there is. Every number in this paper was measured, not estimated.

| DEPENDENCIES                                                              | COMMANDS                                              | VS RIPGREP                                                                          | TESTS                                                           |
|---------------------------------------------------------------------------|-------------------------------------------------------|-------------------------------------------------------------------------------------|-----------------------------------------------------------------|
| <b>0</b><br>pure Rust std; mmap, threads & SIMD via core/libc — no crates | <b>27</b><br>navigate · read · edit · config · verify | <b>1.3×</b><br>faster on a single 380 MB file (~42 vs ~57 ms); parity on many files | <b>168</b><br>+ adversarial review on every feature; ASan-clean |

A FIELD GUIDE TO A SEARCH-AND-EDIT TOOL THAT EARNS AN AGENT'S TRUST

Benchmarks: 10-core Apple Silicon, warm cache, median of 15–21 runs, counting the 99,197 lines containing function in a ~380 MB corpus of real source. Reproduce with `tarn find -c` vs `rg -c` on any large file.

# What's inside

|    |                                           |   |
|----|-------------------------------------------|---|
| 01 | Why agents fight their tools.             | 3 |
| 02 | The surface: one binary, 27 commands.     | 3 |
| 03 | Speed: at parity with ripgrep, zero deps. | 4 |
| 04 | How — mmap, all cores, SIMD, no crates.   | 5 |
| 05 | An honest negative result: the index.     | 6 |
| 06 | The evidence: what ~57,000 calls say.     | 7 |
| 07 | Trust is the product.                     | 8 |

HOW TO READ THIS

## Every claim is cashable.

The performance figures are medians of real runs on a real corpus against real tools (ripgrep 14, ugrep 7.5). Where a thing didn't work, it's in here too — section 05 is a feature we built, measured, and deleted. The tool is on GitHub; `tarn help --json` emits its whole surface for an agent to learn in one call.

# 01

THE PROBLEM

## Why agents fight their tools.

A coding agent drives tools built for humans: open a file, scroll, count lines, hope the line numbers didn't move since the last read. The friction is real and it compounds across a session.

CONTEXT BURNED

### whole files

read into the prompt just to find one symbol. `grep` gives lines without structure; `cat` gives everything without shape.

LINE DRIFT

### off-by-N

an edit keyed to a line number that moved between read and write clobbers the wrong line. No guard, no recovery.

NO MACHINE SURFACE

### prose

human tools answer in prose an agent must re-parse. No stable JSON, no exit-code contract to branch on.

# 02

THE SURFACE

## One binary, 27 commands.

Structural, surgical, and self-describing. `tarn help --json` returns the whole manifest — every command, usage, example, and exit code — so an agent learns the tool in a single call.

NAVIGATE (NO LSP)

### tree • outline • defs • refs • peek • find

Go-to-definition, find-references, repo map, and `grep` that hands you the enclosing definition of every hit — from heuristics, no language server.

EDIT BY MEANING

### replace • insert • delete • apply • patch • batch • rename

`--def <name>` edits a whole function. `--expect` refuses if the target moved. `git diff | tarn patch` applies a unified diff, relocating drifted hunks.

CONFIG & VERIFY

### json • toml • yaml • check • diff

Get/set/delete by path, format-preserving (comments & key order kept). Hygiene gate and a fast line-numbered diff. Exit codes: 0/1/2/3.

THE CONTRACT

## Deterministic, guarded, chainable.


Every edit preserves line endings and never reflows untouched lines; `--dry-run` previews; multi-file `apply/patch` are atomic with rollback; `--json` on read & edit commands lets one command feed the next. An agent can act without a defensive re-read.

# 03 At parity with ripgrep. Zero dependencies.

THE SPEED

ripgrep is the fastest search tool in wide use — SIMD, mmap, a tuned parallel walk. tarn matches it on a directory and *beats* it on a single large file, with not one crate dependency.

COUNT MATCHES IN ONE 380 MB FILE — LOWER IS FASTER  
30,720 matching lines. Warm cache, median of 15 runs (figures bounce ±~15% run to run).



tarn is ~1.3× faster than ripgrep here — and ~45× faster than the system grep (ugrep 7.5: ≈ 2.2 s, off this chart).

COUNT MATCHES ACROSS 3,000 FILES (~380 MB) — LOWER IS FASTER  
The recursive case. Warm cache, median of 15 runs.



Parity — ripgrep edges it ~1.1×. Both fan files out across cores; ripgrep's tuned per-file machinery still leads slightly.

## THE HONEST CEILING

### You can't count faster than you can read.

tarn counts a 380 MB file in ~42 ms; merely *cat*-ing it takes ~36 ms. tarn is at the memory-bandwidth wall — it counts in about the time it takes to read the bytes at all. So "10× faster" on a single scan isn't an optimization, it's physics. (The 10× hides elsewhere — see section 05.)

# 04 Fast with no crates.

THE HOW

ripgrep's raw byte-scan beats anything you'd hand-write in std. tarn wins the difference back three other ways — all standard library, no dependencies.

NO COPY

mmap (libc FFI)

Map the file and scan straight from the page cache — no `fs::read` copy of hundreds of MB. Same FFI `tarn` already uses for `stty`.

ALL CORES

std::thread

ripgrep searches one file on one thread; `tarn` splits it. `\n` is always a UTF-8 boundary, so each chunk counts independently and the sum is exact.


SIMD, STD

NEON via core::arch

A 16-byte-at-a-time `memchr` in NEON intrinsics — which live in `core::arch`, the standard library, not a crate. AddressSanitizer-clean.

THE JOURNEY — TARN FIND -C ON THE SAME 380 MB FILE

Each step is a std-only change. Lower is faster.



563 → ~40 ms: a ~14× speedup, entirely within the standard library.

BONUS: THE DIFF RENDERER

~7 s → ~26 ms on a 40k-line file.

A full-matrix LCS was a latent  $O(n \cdot m)$  memory bomb. Trimming the common prefix/suffix and running the quadratic part only on the differing middle made a one-line change in a 40,000-line file diff in ~26 ms instead of ~7 s — backing every `--diff` flag. And `tarn batch` runs a whole command session in one process: ~10,500 edits/sec (~34× over per-call spawn), since an agent's edit loop is bottlenecked by OS process spawn, not by `tarn`.

# 05 Measure it. Then maybe delete it.

THE CARDINAL RULE

A tool earns trust by what it refuses to ship. We built a persistent navigation index to make repeated defs/refs sublinear — then measured it and pulled it.

THE PREMISE

91%

of a defs query on a 1,688-file repo is re-reading and re-parsing files that didn't change (~142 of ~156 ms). Cache that, the theory went, and queries go ~10×.

THE MEASUREMENT

1.4×

what the built, correct, never-state per-file cache actually delivered. Not 10×.

THE REASON

decode

each tarn call is a fresh process that re-reads and re-deserializes the whole index. The floor is decode + stat, not the parse we cached.

WHAT SHIPPED INSTEAD: NOTHING

## 1.4× doesn't justify persistent state.

A new on-disk format and a new failure surface, for 1.4×, on a tool whose whole appeal is "just point it at files" — not worth it. The prototype was reverted; the analysis is recorded in [docs/INDEX.md](#). A real 10× would need a resident daemon or an mmap'd zero-decode index — a much larger change, parked until the latency is actually felt. **The honest negative result is part of the spec.**


# 06 We stopped guessing. We counted.

## THE EVIDENCE

Every feature so far came from a hunch about how an agent works a repo. Good hunches — but hunches. So we pulled **1,173 real Claude Code session transcripts (a June 2026 snapshot)** and counted what the agents in them actually did: **~57,000 tool calls**, tallied by hand, no sampling. Two findings, two features — neither was on the roadmap before we looked.

### WHAT ~57,000 AGENT TOOL CALLS ACTUALLY WERE

1,173 Claude Code sessions (June 2026 snapshot), every call tallied. ~90% are file/shell ops — tarn's domain.



The rest is thinking and talking. The work an agent does to *code* is overwhelmingly running commands and changing files.

### INSIDE BASH

**34%**  
**grep, by a mile.**

grep was ~15,800 calls — a third of all CLI-utility use, ~100× more than ripgrep. So tarn's search speaks it: `tarn find -e/--regex`, a built-in regex engine, zero crates.

### THE ENGINE

**linear**  
**Thompson/Pike NFA.**

Regex the safe way — no catastrophic backtracking. `(a*)b` on a 50,000-char line returns in 0.00s. Cross-validated against `grep -E`: no unexplained mismatches.

### INSIDE EDIT

**81%**  
**Edits span lines.**

9,661 of 11,872 edits touched more than one line. So `tarn replace <file> A-B <text>` swaps a whole line range for multi-line text — same `--expect / --diff / --dry-run`.

### WHAT IT ISN'T

**A useful subset, named as one.**

The regex engine is a grep-ish subset, not PCRE — no backreferences, lookahead, or `\b`. Unsupported syntax errors *loudly* rather than matching the wrong thing, and literal substring stays the default. Structural `outline/defs` is heuristic, not a parser — multi-line signatures and keyword-less Java/C#/C/C++ methods are now handled by brace-balanced, string/comment-aware scanning; the residual known gap is C# `@"..."` verbatim strings. Built from the data, bounded by the truth — zero crates, 168 tests, every claim cashable.

# 07 Trust is the product.

THE PRODUCT

Speed gets a tool tried; trust gets it adopted. A search tool that returns one wrong-but-plausible answer gets pulled from the harness forever.

**Adversarial review on every feature.** A ruthless reviewer agent gated each change before it shipped — and earned its keep: it caught a real `.env` data-loss bug (a non-UTF-8 byte made `set` wipe the file), a CRLF counting seam, and exit-code contract violations, each fixed at the root.

**Counts proven exact.** The SIMD scanner is differential-tested against a scalar oracle — 900+ fuzz cases plus a 3,000-input run, zero mismatches — and its counts match `rg/grep` on the benchmark corpus. The unsafe NEON code is AddressSanitizer-clean.

**Edits never corrupt.** Line endings (LF/CRLF) and final-newline state are preserved on every edit; format-preserving config edits keep comments and key order; multi-file batches roll back as a unit.

**168 tests, clippy & fmt clean, CI green.** Every benchmark in this paper is reproducible from the repo. Where a number depended on the test machine, it's labelled. Nothing here is aspirational.

## DEFINITION OF DONE

### The tool an agent would choose.

Fast enough that you never reach for `grep`. Structural, so you jump to a symbol instead of scrolling. Guarded, so an edit can't silently land in the wrong place. Self-describing, so any agent learns it in one `tarn help --json`. And honest enough to delete the feature that didn't earn its place. Built by an agent — for the tools it wished it had.



tarn.

THE EDITOR YOUR AGENT WISHES IT HAD · BY AN AGENT, FOR AGENTS